

Image Classification





Multilayer Perceptron

▶ `model.summary()` # Gives the summary of the model

↳ Model: "model"

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	[(None, 150, 150, 3)]	0
conv2d (Conv2D)	(None, 148, 148, 16)	448
max_pooling2d (MaxPooling2D)	(None, 74, 74, 16)	0
conv2d_1 (Conv2D)	(None, 72, 72, 32)	4640
max_pooling2d_1 (MaxPooling2D)	(None, 36, 36, 32)	0
conv2d_2 (Conv2D)	(None, 34, 34, 64)	18496
max_pooling2d_2 (MaxPooling2D)	(None, 17, 17, 64)	0
flatten (Flatten)	(None, 18496)	0
dense (Dense)	(None, 512)	9470464
dropout (Dropout)	(None, 512)	0
dense_1 (Dense)	(None, 2)	1026
=====		
Total params: 9,495,074		
Trainable params: 9,495,074		
Non-trainable params: 0		



```
[ ] history = model.fit_generator(  
    train_generator,  
    steps_per_epoch=100,  
    epochs=10,  
    validation_data=validation_generator,  
    validation_steps=50,  
    verbose=2)
```

```
↳ Epoch 1/10  
Epoch 1/10  
100/100 - 19s - loss: 0.7033 - acc: 0.5220 - val_loss: 0.6771 - val_acc: 0.5750  
Epoch 2/10  
Epoch 1/10  
100/100 - 17s - loss: 0.6766 - acc: 0.5895 - val_loss: 0.6634 - val_acc: 0.5720  
Epoch 3/10  
Epoch 1/10  
100/100 - 17s - loss: 0.6711 - acc: 0.5960 - val_loss: 0.6400 - val_acc: 0.6710  
Epoch 4/10  
Epoch 1/10  
100/100 - 17s - loss: 0.6630 - acc: 0.6135 - val_loss: 0.6347 - val_acc: 0.6630  
Epoch 5/10  
Epoch 1/10  
100/100 - 17s - loss: 0.6368 - acc: 0.6535 - val_loss: 0.6373 - val_acc: 0.6460  
Epoch 6/10  
Epoch 1/10  
100/100 - 17s - loss: 0.6413 - acc: 0.6380 - val_loss: 0.5977 - val_acc: 0.6830  
Epoch 7/10  
Epoch 1/10  
100/100 - 17s - loss: 0.6343 - acc: 0.6530 - val_loss: 0.6089 - val_acc: 0.7030  
Epoch 8/10  
Epoch 1/10  
100/100 - 16s - loss: 0.6299 - acc: 0.6475 - val_loss: 0.5930 - val_acc: 0.7030  
Epoch 9/10  
Epoch 1/10  
100/100 - 17s - loss: 0.6080 - acc: 0.6725 - val_loss: 0.6331 - val_acc: 0.6010  
Epoch 10/10  
Epoch 1/10  
100/100 - 17s - loss: 0.6234 - acc: 0.6615 - val_loss: 0.5843 - val_acc: 0.6960
```

```
[ ] img = tf.keras.utils.get_file('image.jpg', 'https://placekitten.com/200/287')
img = Image.open(img).resize(Input_Image_Shape)
plt.imshow(img)

img = np.array(img)/255.0
print (img.shape)
```

☞ Downloading data from <https://placekitten.com/200/287>
8192/Unknown - 0s 0us/step(150, 150, 3)



We have 2 different classes of Outputs. Cats and Dogs.

```
[ ] plt.imshow(img)
    plt.axis('off')
    _ = plt.title("Prediction: " + Possible_Outputs[Predicted_class])
```




Prediction: Cat






```
plt.figure(figsize=(scale * n_features, scale))  
plt.title(layer_name)  
plt.grid(False)  
plt.imshow(display_grid, aspect='auto', cmap='viridis')
```

🔗 /usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:43: RuntimeWarning: invalid value encountered in true_divide





layers =

25x1 [Layer](#) array with layers:

```
1 'data' Image Input      227x227x3 images with 'zerocenter' normalization
2 'conv1' Convolution      96 11x11x3 convolutions with stride [4 4] and padding [0 0 0 0]
3 'relu1' ReLU             ReLU
4 'norm1' Cross Channel Normalization cross channel normalization with 5 channels per element
5 'pool1' Max Pooling      3x3 max pooling with stride [2 2] and padding [0 0 0 0]
6 'conv2' Grouped Convolution 2 groups of 128 5x5x48 convolutions with stride [1 1] and padding [2 2 2 2]
7 'relu2' ReLU             ReLU
8 'norm2' Cross Channel Normalization cross channel normalization with 5 channels per element
9 'pool2' Max Pooling      3x3 max pooling with stride [2 2] and padding [0 0 0 0]
10 'conv3' Convolution      384 3x3x256 convolutions with stride [1 1] and padding [1 1 1 1]
11 'relu3' ReLU            ReLU
12 'conv4' Grouped Convolution 2 groups of 192 3x3x192 convolutions with stride [1 1] and padding [1 1 1 1]
13 'relu4' ReLU            ReLU
14 'conv5' Grouped Convolution 2 groups of 128 3x3x192 convolutions with stride [1 1] and padding [1 1 1 1]
15 'relu5' ReLU            ReLU
16 'pool5' Max Pooling      3x3 max pooling with stride [2 2] and padding [0 0 0 0]
17 'fc6' Fully Connected    4096 fully connected layer
18 'relu6' ReLU            ReLU
19 'drop6' Dropout          50% dropout
20 'fc7' Fully Connected    4096 fully connected layer
21 'relu7' ReLU            ReLU
22 'drop7' Dropout          50% dropout
23 " Fully Connected        2 fully connected layer
24 'prob' Softmax           softmax
25 " Classification Output  crossentropyex
```

options =

TrainingOptionsSGDM with properties:

Momentum: 0.9000
InitialLearnRate: 0.0015
LearnRateScheduleSettings: [1×1 struct]
L2Regularization: 1.0000e-04
GradientThresholdMethod: 'l2norm'
GradientThreshold: Inf
MaxEpochs: 10
MiniBatchSize: 50
Verbose: 1
VerboseFrequency: 50
ValidationData: []
ValidationFrequency: 50
ValidationPatience: Inf
Shuffle: 'once'
CheckpointPath: ''
ExecutionEnvironment: 'gpu'
WorkerLoad: []
OutputFcn: []
Plots: 'none'
SequenceLength: 'longest'
SequencePaddingValue: 0
DispatchInBackground: 0

Initializing input data normalization.

```
=====|
| Epoch | Iteration | Time Elapsed | Mini-batch | Mini-batch | Base Learning |
|       |           | (hh:mm:ss)  | Accuracy  | Loss       | Rate         |
|=====|
| 1 | 1 | 00:00:01 | 52.00% | 1.8875 | 0.0015 |
| 2 | 50 | 00:00:32 | 92.00% | 0.1245 | 0.0015 |
| 3 | 100 | 00:01:04 | 94.00% | 0.2004 | 0.0015 |
| 4 | 150 | 00:01:36 | 100.00% | 0.0031 | 0.0015 |
| 5 | 200 | 00:02:08 | 100.00% | 0.0049 | 0.0015 |
| 6 | 250 | 00:02:40 | 100.00% | 0.0028 | 0.0015 |
| 7 | 300 | 00:03:12 | 100.00% | 0.0023 | 0.0015 |
| 8 | 350 | 00:03:44 | 100.00% | 0.0004 | 0.0015 |
| 9 | 400 | 00:04:16 | 100.00% | 0.0026 | 0.0015 |
| 10 | 450 | 00:04:48 | 100.00% | 0.0002 | 0.0015 |
| 10 | 480 | 00:05:08 | 100.00% | 0.0012 | 0.0015 |
|=====|
```

numCorrect =

585

fracCorrect =

0.9750




```
numCorrect =
```

585

```
fracCorrect =
```

0.9750

```
ans =
```

[ConfusionMatrixChart](#) with properties:

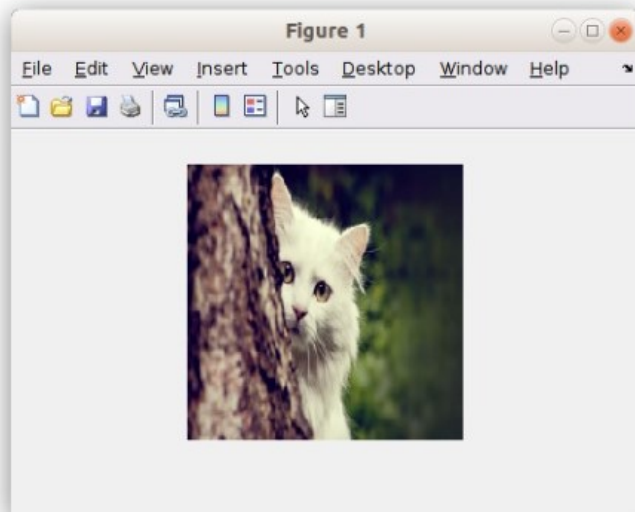
NormalizedValues: [2×2 double]
ClassLabels: [2×1 categorical]

Show [all properties](#)

```
Predict =
```

[categorical](#)

Cat



```
>> img = imread("Dog.jpg");  
img = imresize(img,[227 227]);  
imshow(img)  
Predict = classify(Petnet,img)
```

```
Predict =
```

[categorical](#)

Dog

```
>>
```

